



# HYPERSCAN UPDATES



[WWW.HYPERSCAN.IO](http://WWW.HYPERSCAN.IO)

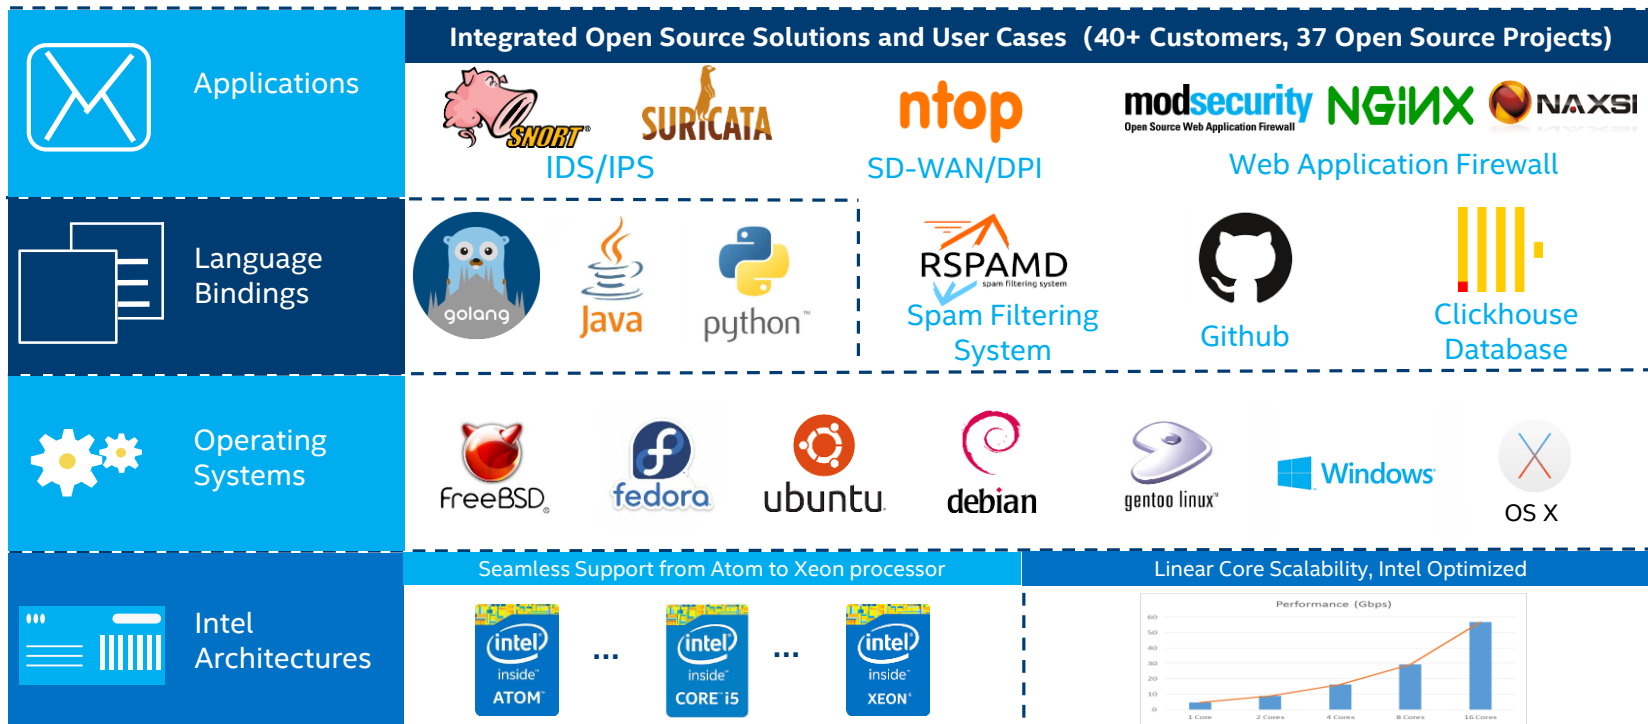
[xiang.w.wang@intel.com](mailto:xiang.w.wang@intel.com)

# Agenda

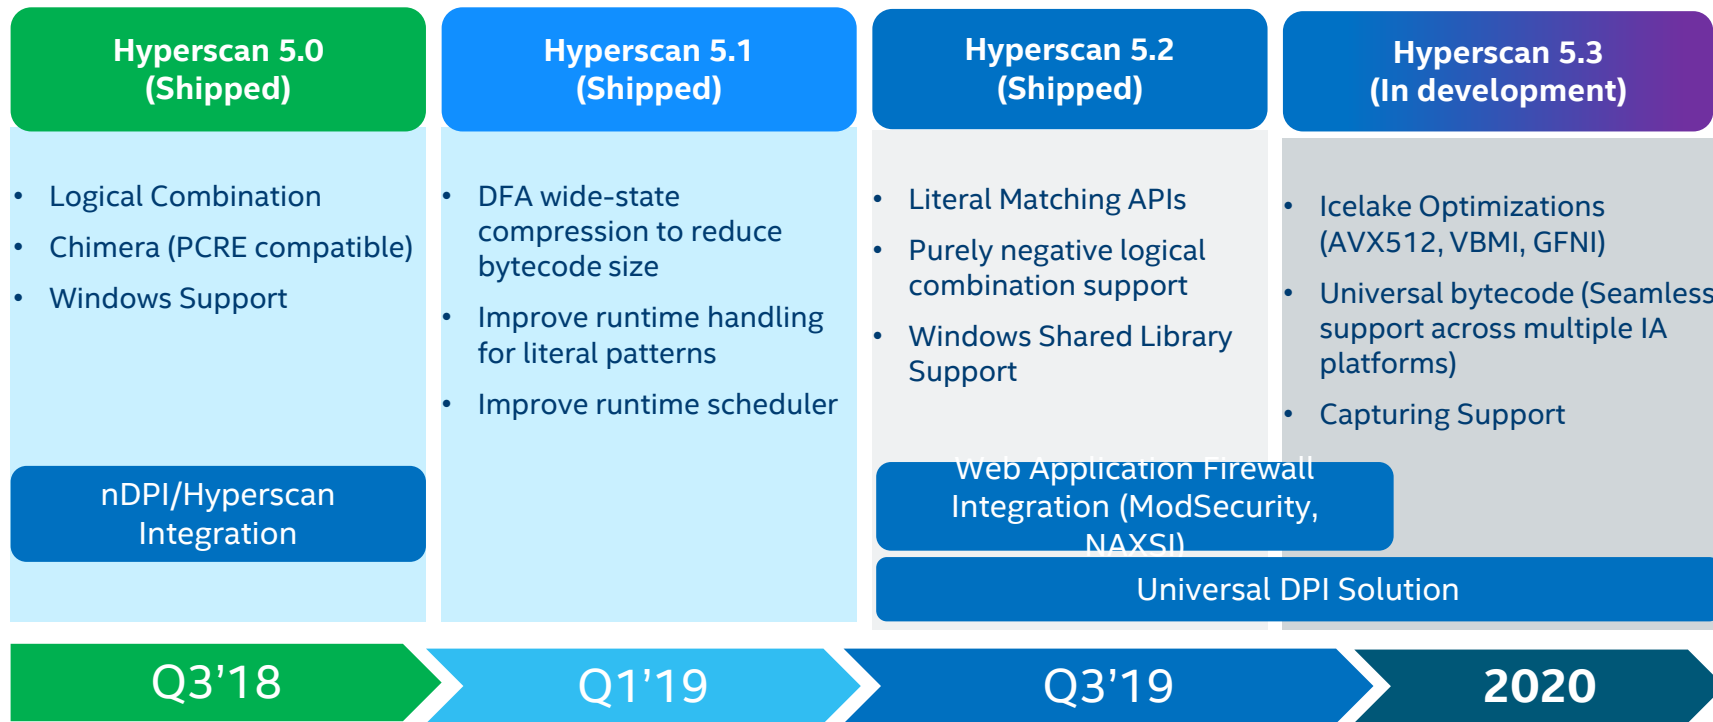
- Hyperscan Ecosystem Overview
- Hyperscan Roadmap
- Hyperscan Performance
- nDPI Case Study
- New open source DPI on FD.io: Universal DPI

# Hyperscan:

An industry fastest Regular Expression, Literal Matching Algorithm on Intel platform



# Hyperscan Roadmap

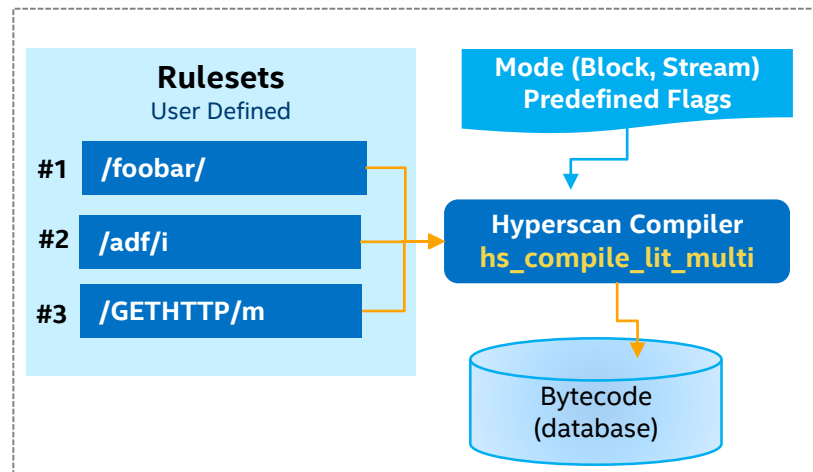


\* Hyperscan is supported on all Intel Platform from Atom, Core to Xeon Scalable Processor Family

# Release 5.2 (Q3'19)

## Literal Matching APIs

- A new set of APIs **hs\_compile\_lit()** and **hs\_compile\_lit\_multi()** for literal matching.
- No need to convert patterns into HEX representation to avoid regex syntaxes.
- Simplified and faster compile and runtime process.



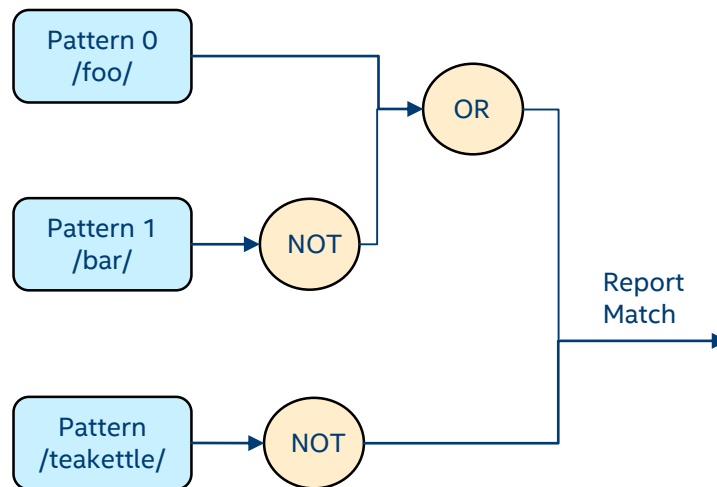
# Release 5.2 (Q3'19)

## Logical combination

- Report matches only when find defined logical combination (unordered AND, NOT and OR) of patterns
- Add support for **purely negative** patterns, e.g. (pattern0 OR (NOT pattern1)) AND (NOT pattern2)

## Windows

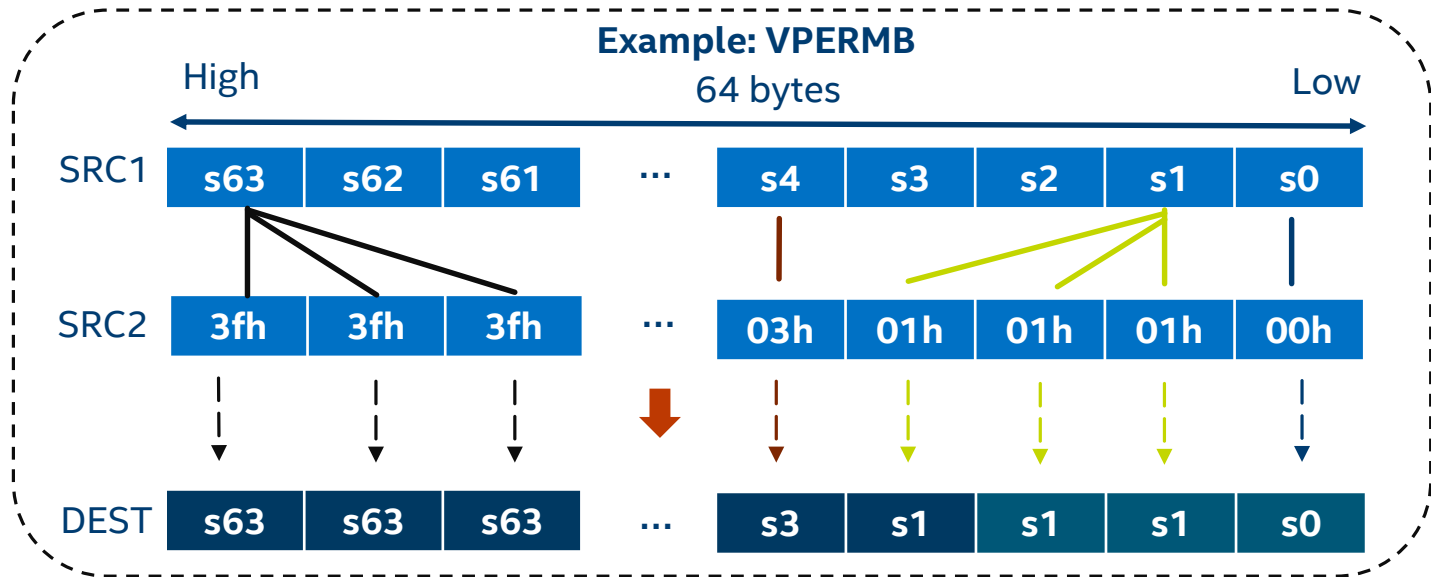
- Add support for DLL shared library build for Window 32/64-Bit OS



# Release 5.3 (In Development)

## Icelake optimizations

- AVX512 VBMI utilization for literal matching, acceleration models, lookahead, etc.



# Release 5.3 (In Development)

## Capturing Support

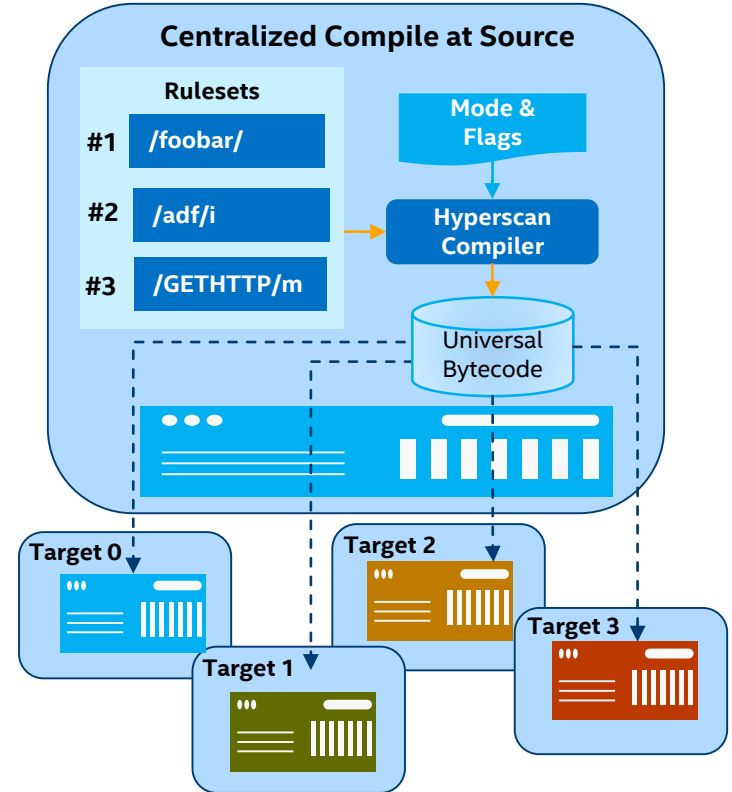
- Hyperscan ignores capturing syntax and only reports matching offsets for the whole pattern
- Add capturing syntax support widely used in WAF, IDS/IPS, etc  
e.g. foobar(badge)brush, offsets of “badge” will be reported together with overall pattern matching offsets



# Release 5.3 (In Development)

## Universal Bytecode

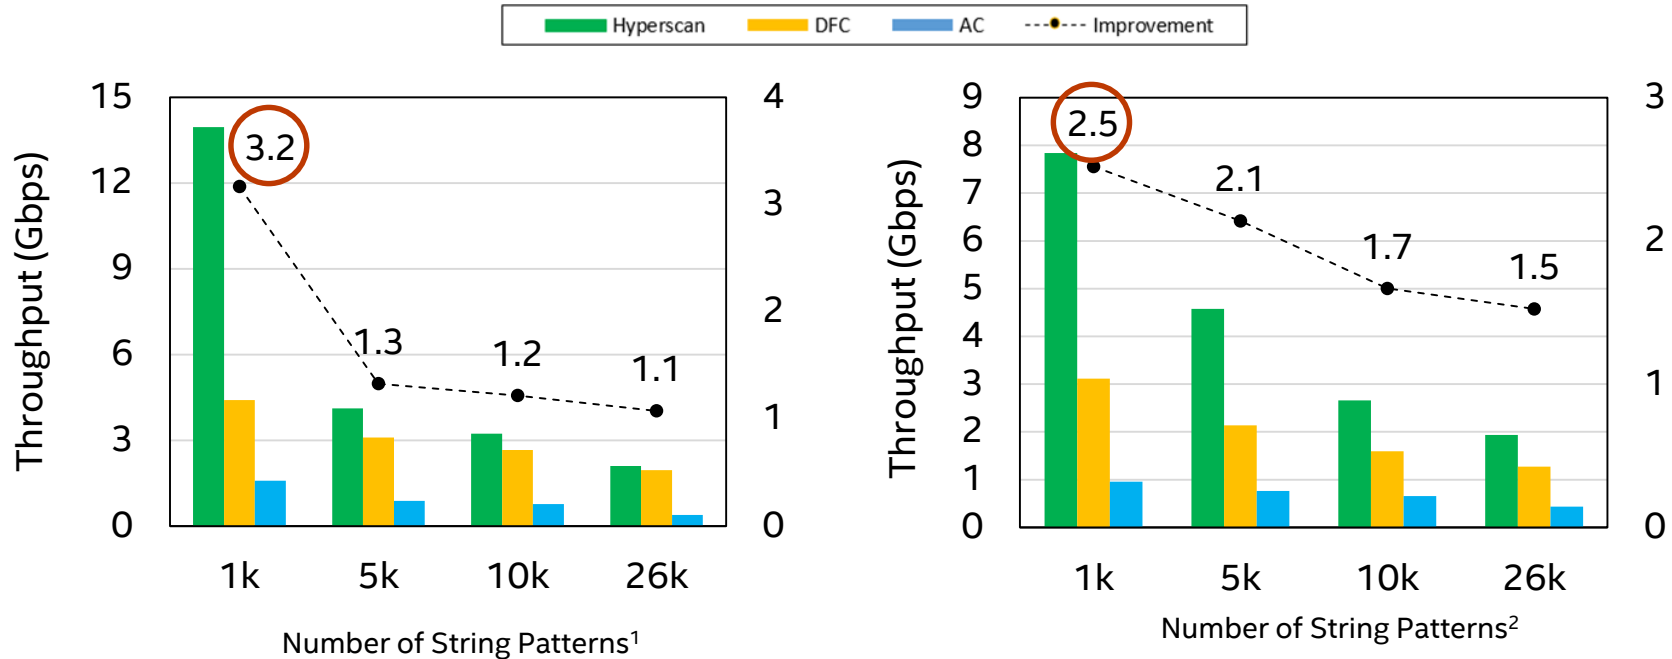
- **Centralized compile** of a universal bytecode at source machine
- Distribute this universal bytecode to multiple target machines with **different Intel architectures**
- Each target machine automatically pick up bytecode section best for its architecture during scan



# Evaluation of Hyperscan

- Performance of **literal matching** and **regex matching** vs. state-of-the-art solutions
- Experiment setup:
  - Machine: Intel Xeon Platinum 8180 CPU @ 2.50GHz (48 GB of RAM)
    - ❖ Runs with a single core
    - ❖ GCC 5.4
  - Ruleset: Snort Talos (May 2015), Snort ET-Open 2.9.0, Suricata rulesets 4.0.4
  - Workload: random traffic, real-world web traffic

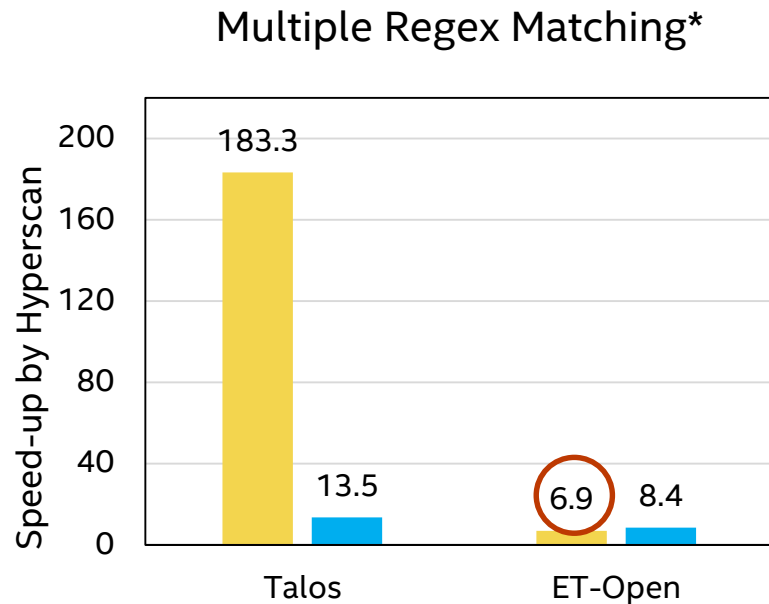
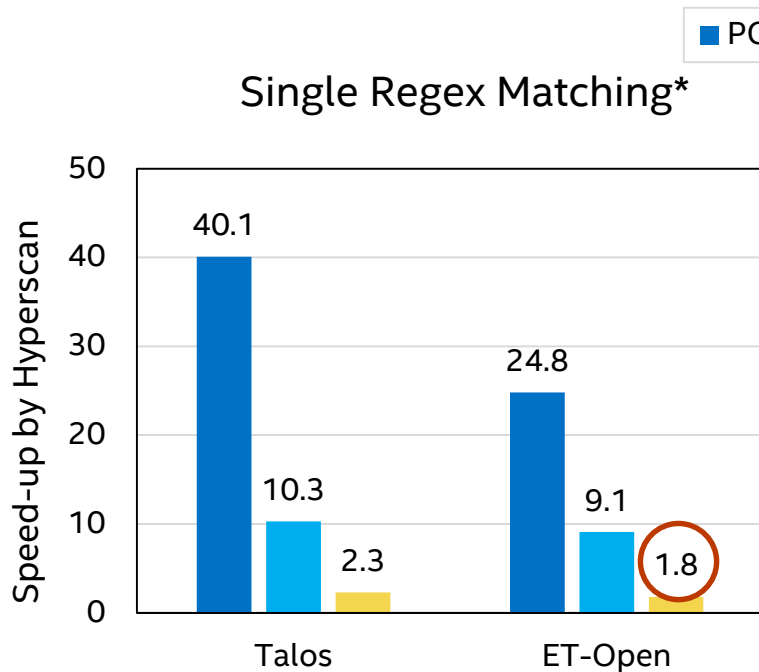
# Multi-literal Matching Performance with Snort ET-Open



<sup>1</sup> Random workload.

<sup>2</sup> Real web traffic trace.

# Regex Matching Performance



\* Test with Snort Talos (1,300 regexes) and ET-Open (2,800 regexes) rulesets under real Web traffic trace.

# nDPI

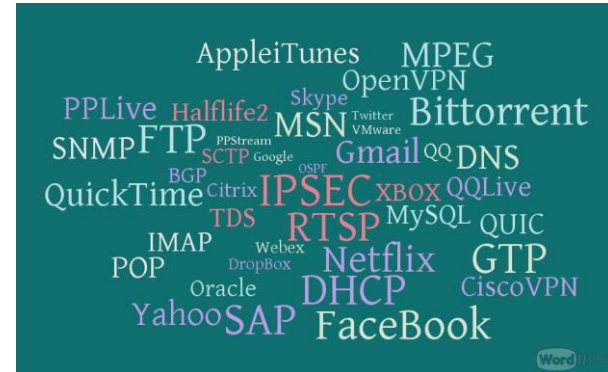
Open source DPI solution under GNU LGPL license  
(<https://github.com/ntop/nDPI>)



Support more than 240 protocols/applications

- Messaging (Facebook, Whatsapp)
- Multimedia (YouTube, iTunes)
- Conferencing (Webex, CitrixOnLine)
- Streaming (iQiyi, Netflix)
- Business (VNC, Citrix)

Flexible to add new protocol dissector  
Hyperscan Accelerated literal matching



# nDPI Database

- Maintain a protocol database that mostly contains HTTP URLs
- Multiple literal matching (Hyperscan or Aho Corasick) to detect URLs in packets
- Example: Wechat rules

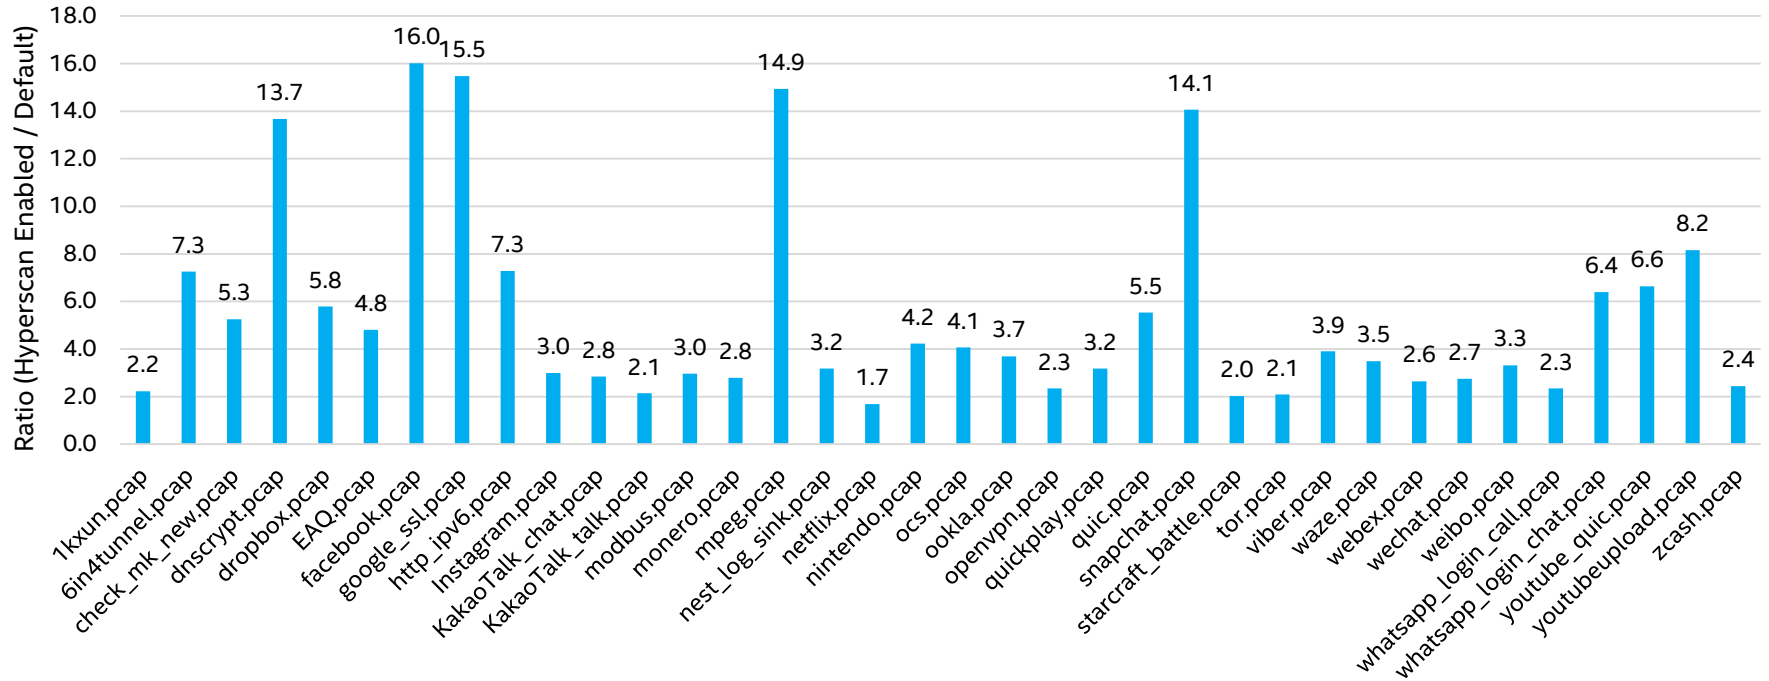
```
{ ".wechat.com", NULL, "\\..wechat\\.com", "WeChat", NDPI_PROTOCOL_WECHAT,  
NDPI_PROTOCOL_CATEGORY_CHAT, NDPI_PROTOCOL_FUN },  
  
{ ".wechat.org", NULL, "\\..wechat\\.org", "WeChat", NDPI_PROTOCOL_WECHAT,  
NDPI_PROTOCOL_CATEGORY_CHAT, NDPI_PROTOCOL_FUN },  
  
{ ".wechatapp.com", NULL, "\\..wechatapp", "WeChat", NDPI_PROTOCOL_WECHAT,  
NDPI_PROTOCOL_CATEGORY_CHAT, NDPI_PROTOCOL_FUN },  
  
{ ".we.chat", NULL, "\\..we\\.chat", "WeChat", NDPI_PROTOCOL_WECHAT,  
NDPI_PROTOCOL_CATEGORY_CHAT, NDPI_PROTOCOL_FUN },
```

# nDPI Performance Evaluation

- nDPI® uses an application called ndpireader to give an application level protocol analysis on each pcap.
- Throughput is measured in Kpps(packets/sec).
  - The nDPI® throughput, Average Memory and Setup Time are noted for each run.
- nDPI® reads a total of 76 pcaps acquired from top sites. In this report we are only showing performance of top 36 pcaps.

# nDPI Performance Ratio – Hyperscan vs Aho Corasick

Performance Gain Hyperscan Enabled vs Default nDPI®



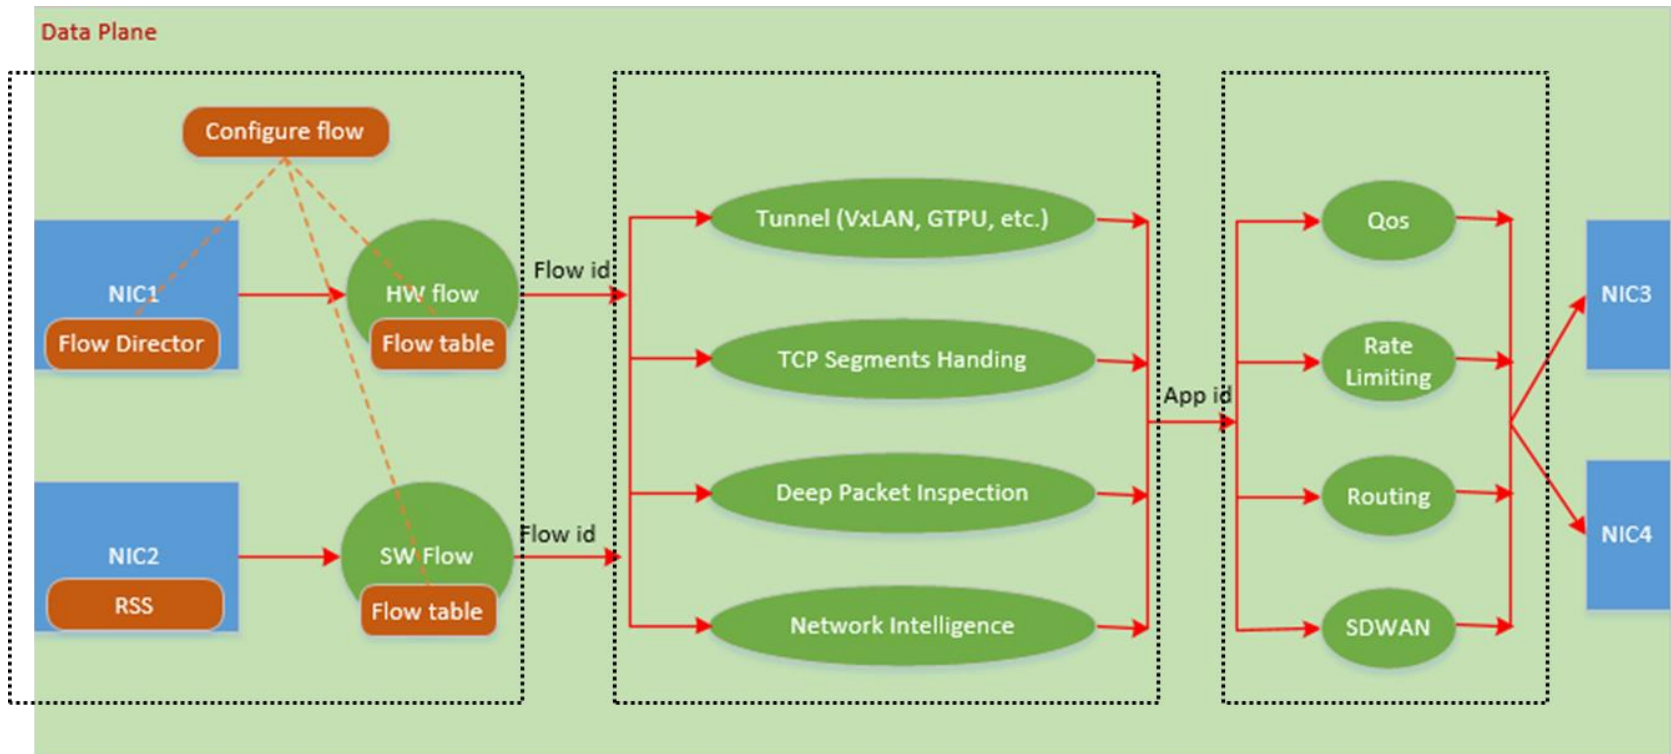
Pcap Based/ 1C1T/ nDPI reader, Intel® Xeon® Gold 6152 @ 2.10GHz



# UDPI: Universal Deep Packet Inspection

- A new project under FD.io: <https://wiki.fd.io/view/UDPI>
- Flow Classification & Expiration
  - HW flow offloading leveraging rte\_flow on DPDK
  - SW flow classification
- Application Detection
  - Leverage Hyperscan Stream Mode
  - Reassembly TCP segments on the fly
- Application-based Actions
  - HQos, Rate Limiting, Policy Routing, SDWAN, etc.
- Support Hundreds of Protocols & Applications
  - TLS/HTTPS, HTTP, DNS, QUIC, etc.

# UDPI Architecture



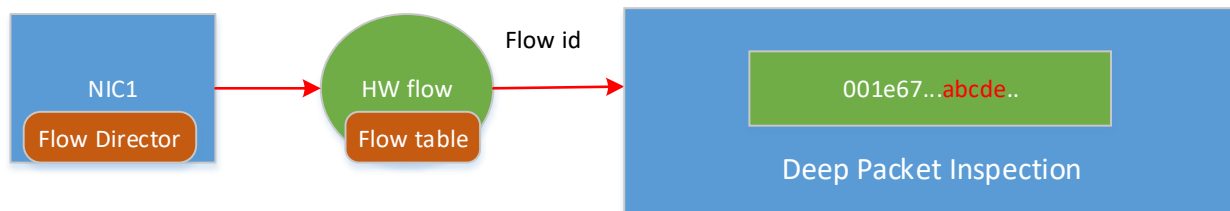
1. Flow Classification

2. Application Detection

3. Application-based Action

# Block Mode and TCP Segments Reassembly

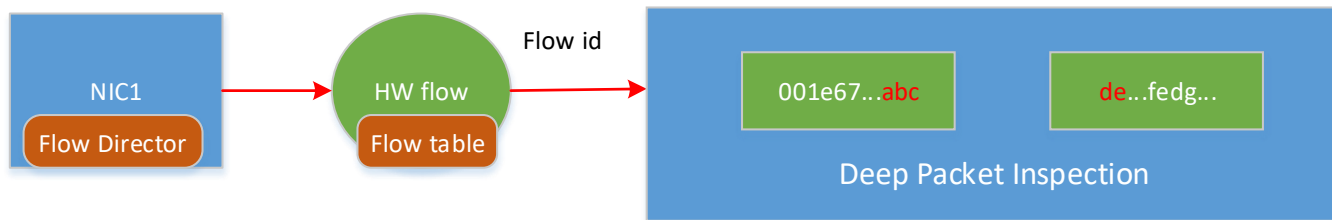
- Block mode can scan rules only in a complete payload.
- If defining a rule "abcde", then for Block Mode, "abcde" should be in a complete PDU payload.



- Requirements for TCP Segments process:
  - Reassembly TCP segments first to a complete PDU payload.
  - Scan PDU payload through Block mode.
  - Fragment TCP segments again.
- This degrades the performance.
- Most DPI open source projects leveraging hyperscan performs in this way.

# Stream Mode and TCP Segments Reassembly

- Stream mode can scan rules straddling into different TCP segments.
- If defining a rule "abcde", then for Stream Mode, then "abc" can be reside in packet 1, and "de" can be in packet 2.



- Requirements for TCP Segments process:
  - Reassemble TCP segments reassembly on the fly.
  - Can handle out-of-order tcp segments.
  - Can handle overlapping segments.
- This helps to improve the performance.
- UDPI project is implemented in this way.

# UDPI Founders and Committers

13 organizations joined and 20 initial committers.



ZTE



inspur



Tencent 腾讯



## Committers:

- Xiang Wang [@ Intel](#),
- Yang Hong [@ Intel](#),
- Harry Chang [@ Intel](#),
- Jian Gu [@ ZTE](#),
- Jianghua Shan [@ China Telecom](#),
- Yang Zhang [@ China Telecom](#),
- Xingfu Li [@ HuachenTel](#),
- Shuai Wu [@ Inspur](#),
- Yuying Xia [@ Yxlink](#),
- Chenggang Fan [@ Sunyainfo](#),
- Feng Gao [@ Tencent](#),
- Zhong Liu [@ China Unicom](#),
- Yong Zhao [@ Huawei](#),
- Haiquan Chen [@ QingCloud](#),
- Jim Thompson [@ Netgate](#),
- Pengjie Li [@ Alibaba](#),
- Zhao Zhang [@ Alibaba](#),
- Zhangpeng Xie [@ Alibaba](#),
- Drenfong Wang [@ Intel](#),
- Hongjun Ni [@ Intel](#),

# Conclusion

- Solid and mature with better performance than state-of-the-art solutions
- Still a WIP in many senses, especially for Icelake optimizations
- Used in large number of commercial deployments and delivers significant performance boost for DPI applications
- UDPI provides a complete open source DPI solution and call for collaboration and contributions