

PACKET PROCESSING PERFORMANCE & POWER EFFICIENCY

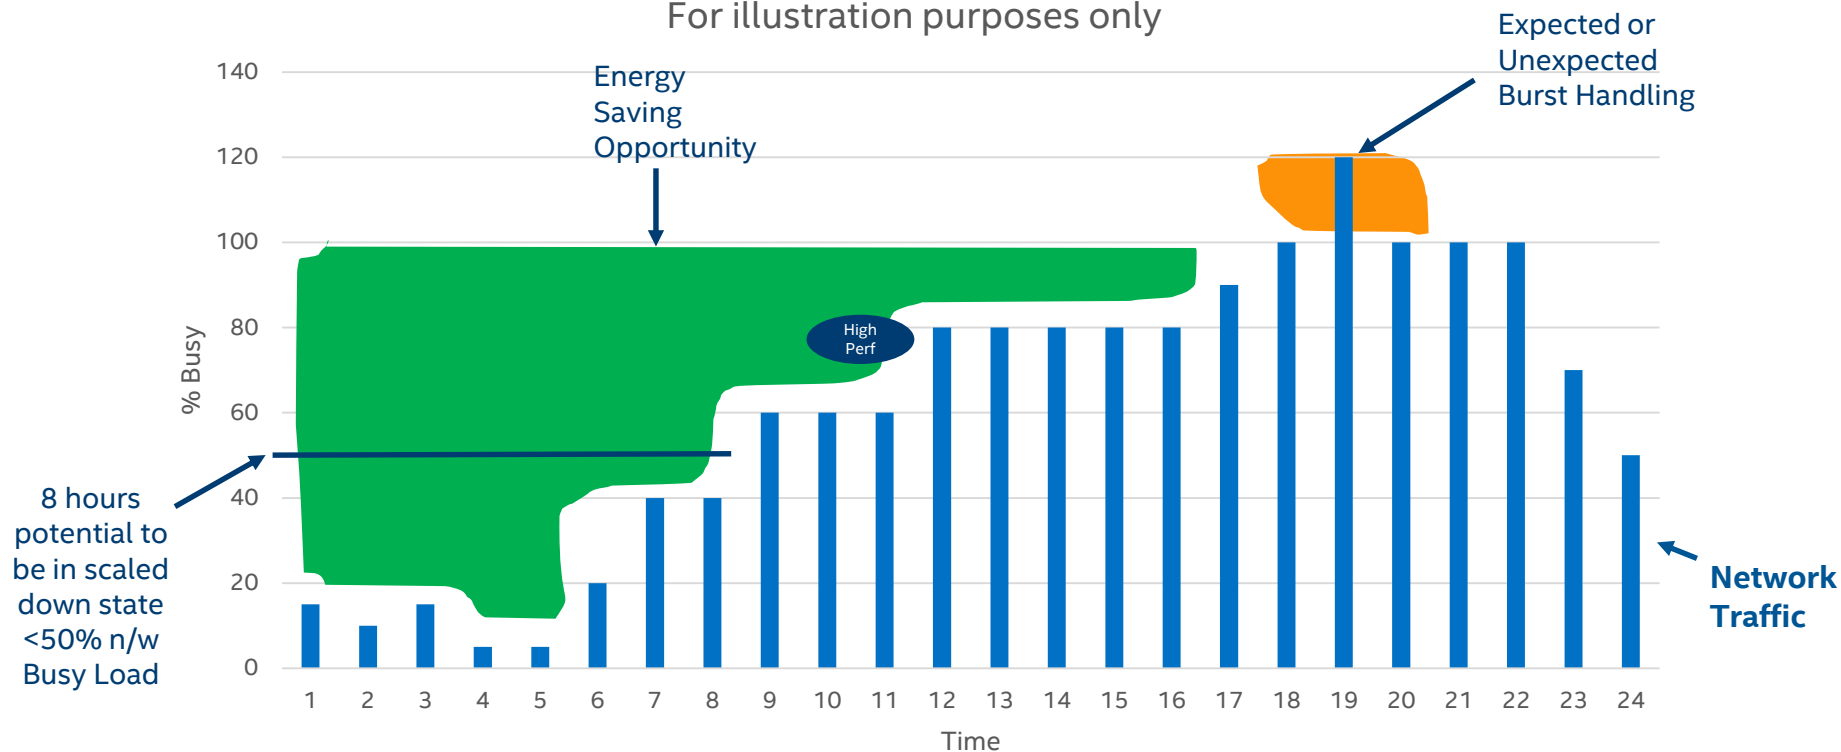
Ma Liang

Senior Software Engineer INTEL

07, Sep, 2019

Mapping Power Usage to Network Traffic

For illustration purposes only

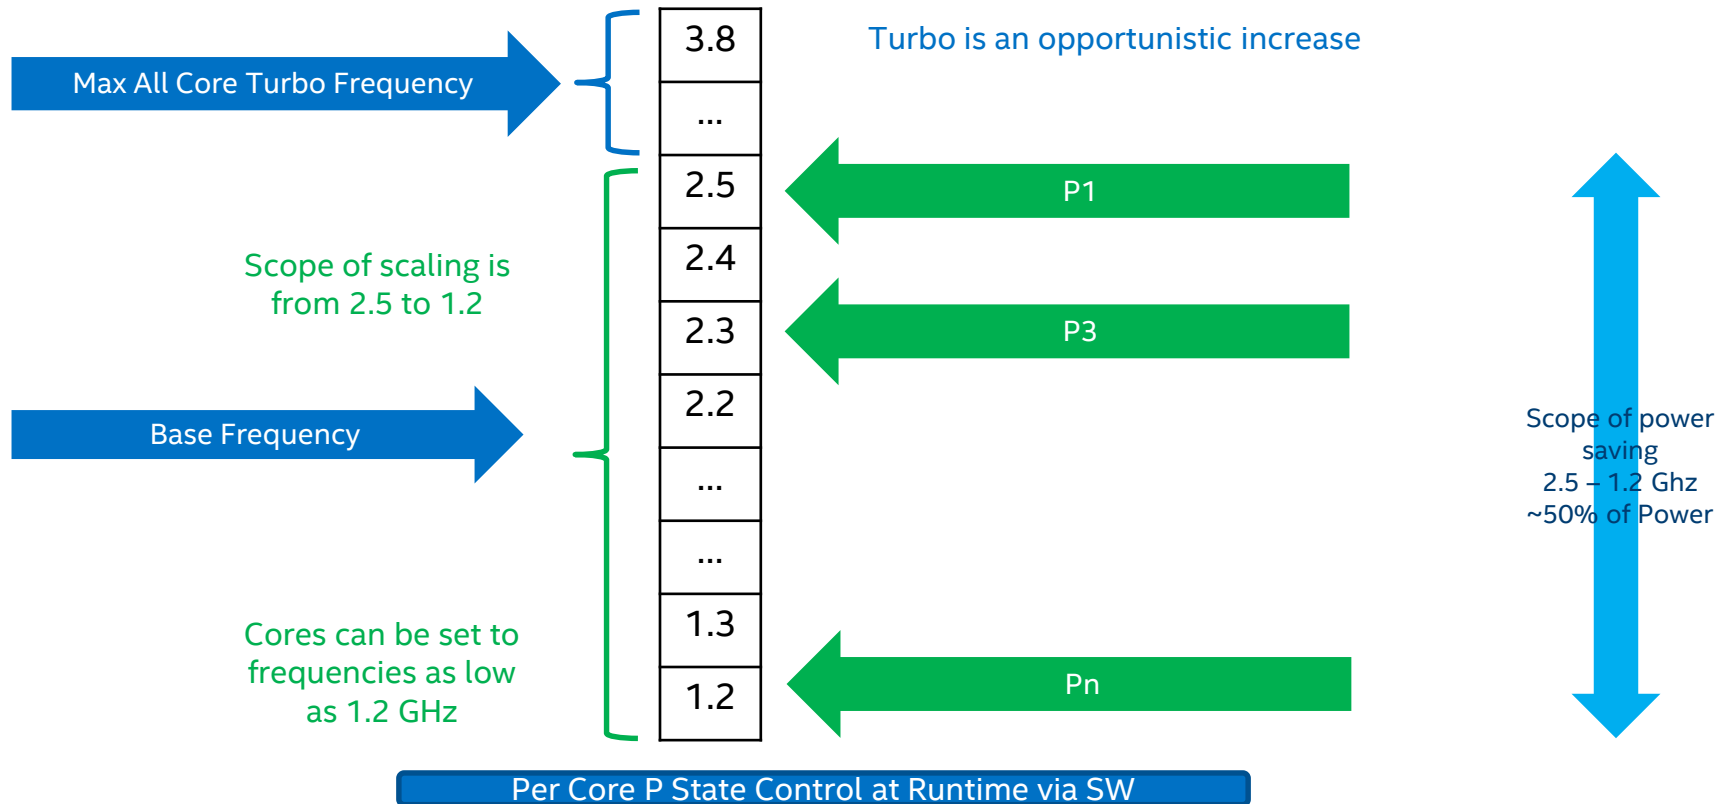


Summary Table of Power Management Features

| Feature | Summary |
|--------------------------------|--|
| P-state | Changing frequency, works per core. Execution continues. |
| C-state | Turning off execution of cores and instructions. Fast Exit(C1) and Longer Exit (C6), power saving versus exit latency. |
| Turbo Boost | Allows for exceeding base frequencies, opportunistic frequency increase. Can be controlled per core. |
| Uncore Frequency Scaling (UFC) | Interconnect and L3 shared cache frequency scaling for energy efficiency |
| Hardware P-State | Intel® Speed Select Technology (Hardware P-state, HWP) is a capability for cooperative hardware + software performance control |
| Intel® SST-Base Frequency | Intel® Speed Select Technology (Base Frequency) is a capability which allow application to choose High Priority or Standard Priority cores |


P-States Overview

Frequency range in GHz of Scalable 8180



Intel® SST-Base Frequency


- Trade off base frequency between higher and lower priority cores
- Improve overall performance by giving critical cores higher frequency
- Intel® Speed Select Technology - Base Frequency (Intel® SST-BF) is enabled in BIOS at boot time; activated/deactivated at run time



All SKUs, frequencies, and performance estimates are PRELIMINARY and can change without notice.


Intel® SST-Core Power

- TDP is Fixed
- Distribute surplus power to cores based on SW assigned weights/priorities
- Improve performance by directing frequency to high priority cores
- Best Effort



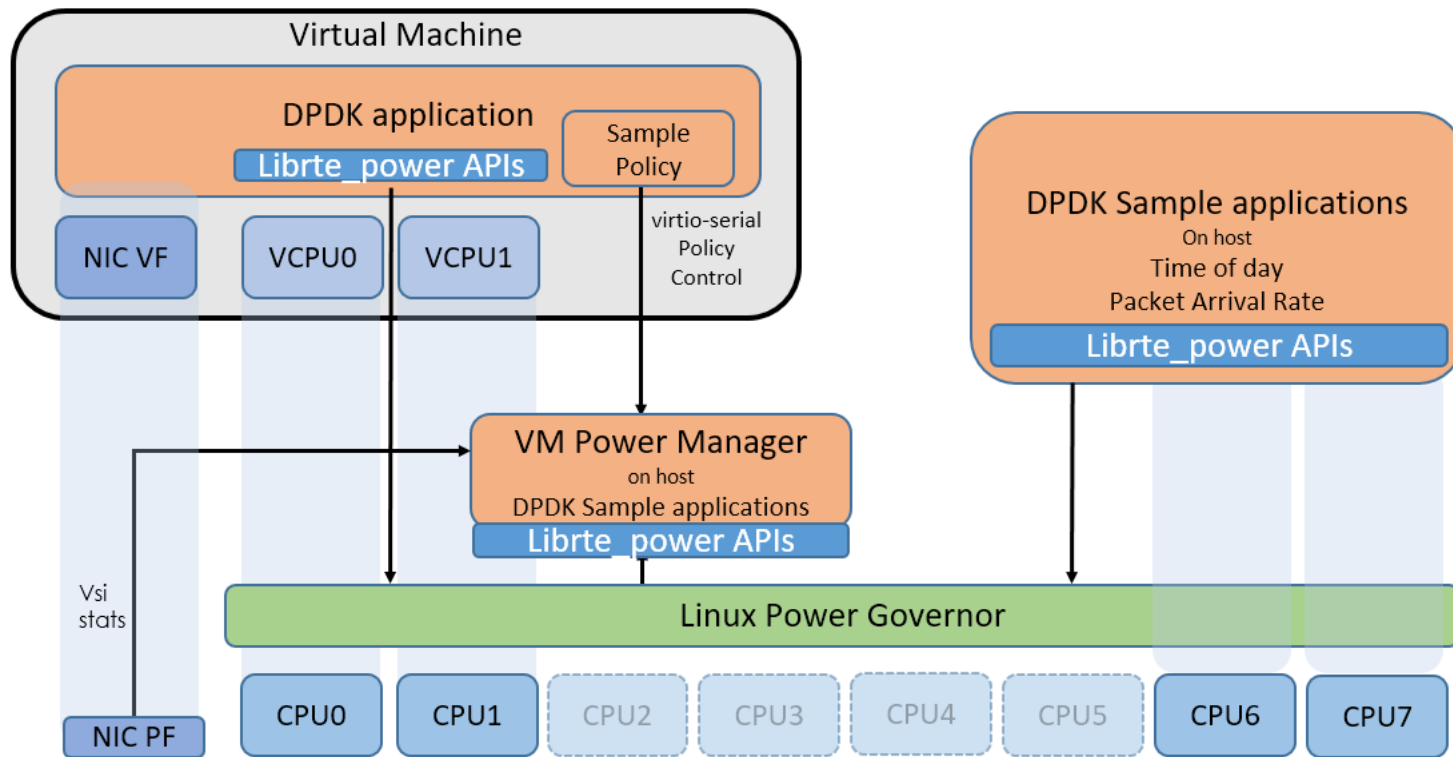
All SKUs, frequencies, and performance estimates are PRELIMINARY and can change without notice.

Performance Gain Potential



Many use cases: boosting 1-2 cores will benefit workloads

Existing DPDK Power Capabilities



DPDK API Reference

rte_power.h File Reference

```
#include <rte_common.h>
#include <rte_byteorder.h>
#include <rte_log.h>
#include <rte_string_fns.h>
```

Go to the source code of this file.

Data Structures

struct [rte_power_core_capabilities](#)

Typedefs

```
typedef uint32_t(* rte\_power\_freqs\_t)(unsigned int lcore_id, uint32_t *freqs, uint32_t num)
typedef uint32_t(* rte\_power\_get\_freq\_t)(unsigned int lcore_id)
typedef int(* rte\_power\_set\_freq\_t)(unsigned int lcore_id, uint32_t index)
typedef int(* rte\_power\_freq\_change\_t)(unsigned int lcore_id)
typedef int(* rte\_power\_get\_capabilities\_t)(unsigned int lcore_id, struct rte\_power\_core\_capabilities *caps)
```

Functions

```
int rte\_power\_set\_env (enum power_management_env env)
void rte\_power\_unset\_env (void)
enum power_management_env rte\_power\_get\_env (void)
int rte\_power\_init (unsigned int lcore_id)
int rte\_power\_exit (unsigned int lcore_id)
```

Variables

```
rte\_power\_freq\_change\_t rte\_power\_freq\_up
rte\_power\_freq\_change\_t rte\_power\_freq\_down
rte\_power\_freq\_change\_t rte\_power\_freq\_max
rte\_power\_freq\_change\_t rte\_power\_freq\_min
rte\_power\_freq\_change\_t rte\_power\_turbo\_status
rte\_power\_freq\_change\_t rte\_power\_freq\_enable\_turbo
rte\_power\_freq\_change\_t rte\_power\_freq\_disable\_turbo
```

DETERMINING AND PREDICTING LOAD

Meeting the needs of an on demand network

Scale always on DPDK performance with the network demand

Common Challenges

- Always On
 - Adjust PMD cores frequency to adjust to packet demand
 - Potential to save power drawn per core using frequency scaling
 - ++ from sleeping


Speed of Re(Action)

- Challenge: Fast Scale Up to react to increases in n/w traffic
- Time = queueing/buffering

Challenge: Fast Monitor & Reaction Time

- Closer to hardware gives faster reaction time


Move to Policy based control



Apply Power Where and When it's needed

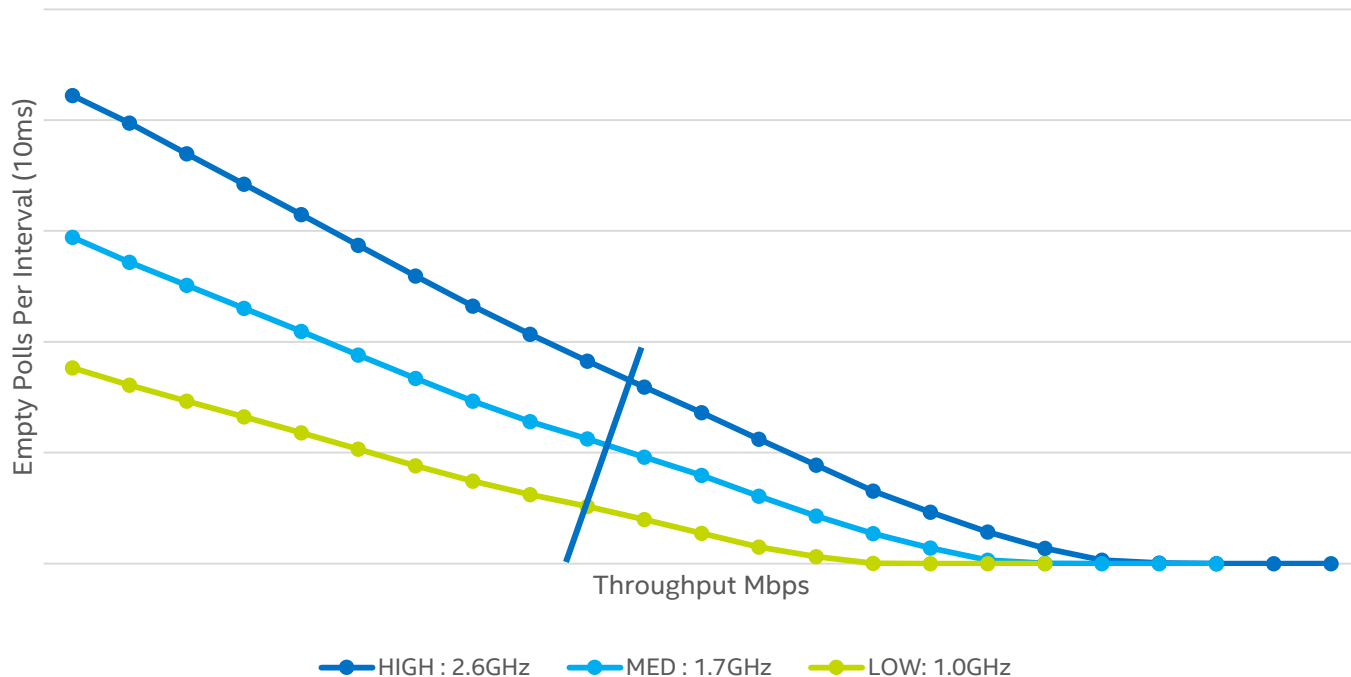
Poll Loop Work Rate Detection (PMD Load%)

- CPU Load is always 100% for DPDK PMD Poll Loops
- Actual workload may be zero (processing zero packets)
- Use PMU counters to calculate the actual work done
- Use the ratio between Branch Hits and Branch Misses
- Ratio is low when tight code loop (empty polling), and significantly is higher when processing packets (due to larger code path)
- Almost linear with traffic rate




Empty Poll Number Driven Model

Empty Polls Per Interval vs. Throughput (@ 3 CPU Frequencies)



Latency Comparison

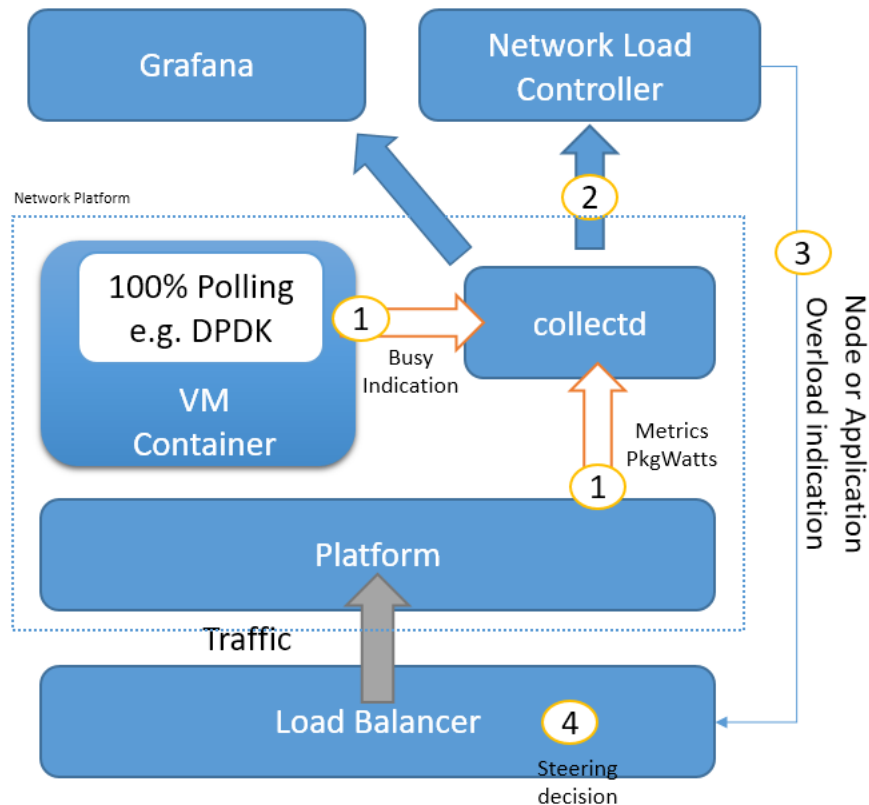


ECO-SYSTEM ENABLEMENT

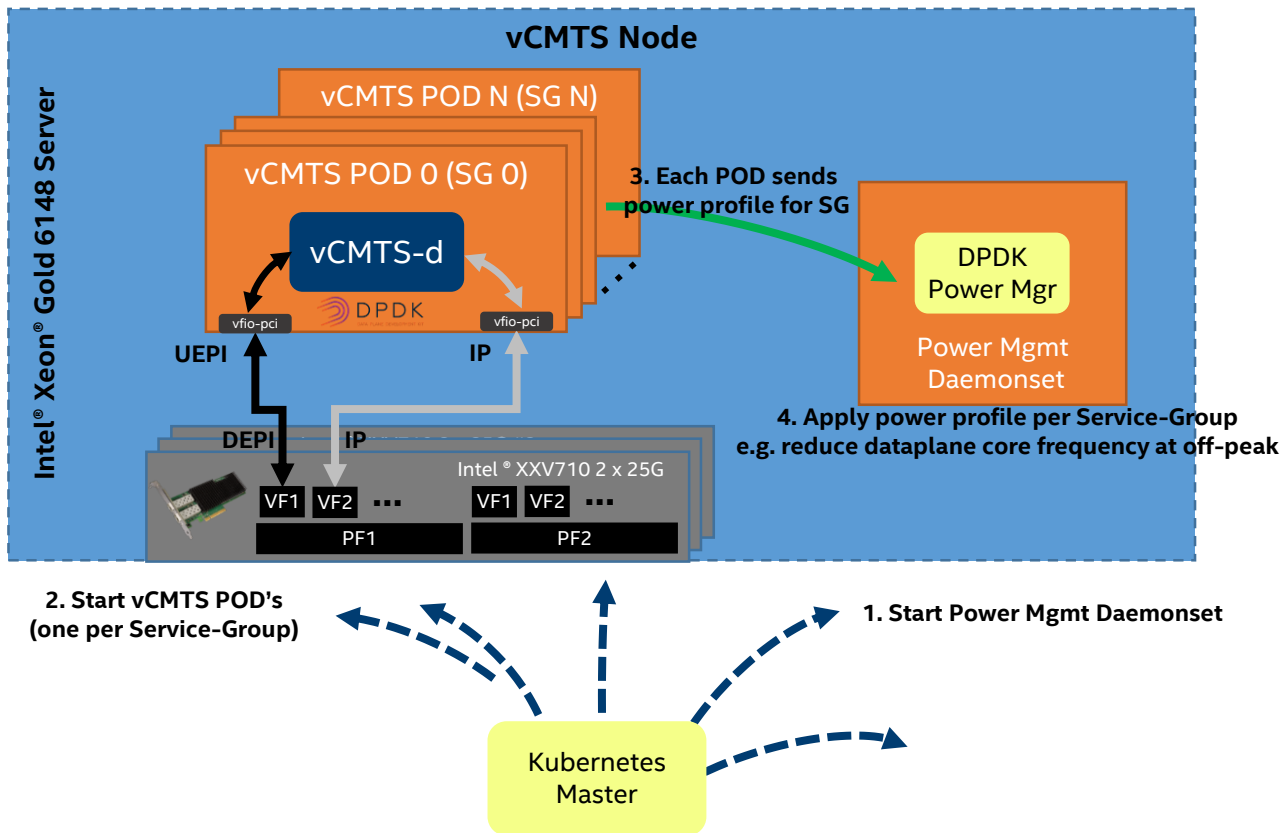
Telemetry Integration

- Metric **Busy Indication**
- Metric **PkgWatts**
- collectd (next) with new dpdk_plugin, updated platform power metrics
- DPDK 19.08 with new telemetry mode sample

| Metric #1 Busy Indication | Metric #2 PkgWatts | Action |
|------------------------------|-----------------------|--------------|
| No | No | Steady State |
| Yes | No | Backoff |
| No | Yes | Backoff |



Kubernetes Integration Example



Software Reference

- DPDK APIs available
 - Application APIs to support in band and out of band use cases
 - http://dpdk.org/doc/api/rte__power_8h.html
- Sample applications
 - L3fwd-power
- Presentations
 - <https://dpdksummit.com/Archive/pdf/2017Userspace/DPDK-Userspace2017-Day2-8-Power.pdf>
 - https://www.dpdk.org/wp-content/uploads/sites/35/2018/10/pm-01-DPDK_Summit18_PowerManagement.pdf
- Power Tools Repo Info
 - <https://github.com/intel/commspovermanagement>

